

DM505 Database Design and Programming

DM576 Database Systems

Panagiotis Tampakis

ptampakis@imada.sdu.dk

Real SQL Programming

SQL in Real Programs

- We have seen only how SQL is used at the generic query interface – an environment where we sit at a terminal and ask queries of a database
- Reality is almost always different: conventional programs interacting with SQL

Options

1. Code in a specialized language is stored in the database itself (e.g., PSM, PL/pgsql)
2. SQL statements are embedded in a *host language* (e.g., Java)
3. Connection tools are used to allow a conventional language to access a database (e.g., CLI, JDBC, psycopg2)

Stored Procedures

- PSM, or “*persistently stored modules*,” allows us to store procedures as database schema elements
- PSM = a mixture of conventional statements (if, while, etc.) and SQL
- Let us do things we cannot do in SQL alone

Procedures in PostgreSQL

```
CREATE PROCEDURE <name>  
  ([<arguments>]) AS $$  
  <program>$$ LANGUAGE <lang>;
```

- PostgreSQL only supports functions:

```
CREATE FUNCTION <name>  
  ([<arguments>]) RETURNS VOID AS $$  
  <program>$$ LANGUAGE <lang>;
```

Procedu
now su

Parameters for Procedures

- Unlike the usual name-type pairs in languages like Java, procedures use mode-name-type triples, where the *mode* can be:
 - IN = function uses value, does not change
 - OUT = function changes, does not use
 - INOUT = both

Function RETURN Types

- VOID
- Traditional data type (e.g. INT, VARCHAR etc.)
- Composite Types (e.g. ROW and RECORD)
- SETOF (e.g ROWS/RECORDS)
- TABLE
- TRIGGER
- Etc.

Example: Stored Procedure

- Let's write a procedure that takes two arguments b and p , and adds a tuple to **Sells(bar, beer, price)** that has bar = 'C.Ch.', beer = b , and price = p
 - Used by Cafe Chino to add to their menu more easily

The Procedure

```
CREATE PROCEDURE ChinoMenu (
```

```
  IN b    CHAR(20),  
  IN p    REAL
```

Parameters are both
read-only, not changed

```
) AS $$
```

```
INSERT INTO Sells  
VALUES(' C.Ch.' , b, p);
```

The body ---
a single insertion

```
$$ LANGUAGE plpgsql;
```

Invoking Procedures

- Use SQL/PSM statement CALL, with the name of the desired procedure and arguments

- **Example:**

```
CALL ChinoMenu('Eventyr', 50);
```

- Functions used in SQL expressions wherever a value of their return type is appropriate
- PostgreSQL: CALL → SELECT

```
SELECT ChinoMenu('Eventyr', 50);
```

Kinds of PL/pgsql statements

- **Return statement:** RETURN <expression> returns value of a function
 - Like in Java, RETURN terminates the function execution
- **Declare block:** DECLARE <name> <type> used to declare local variables
- **Groups of Statements:** BEGIN . . . END
 - Separate statements by semicolons

Kinds of PL/pgsql statements

- **Assignment statements:**
 <variable> := <expression>;
 - Example: b := 'Od.Cl.';
- **Statement labels:** give a statement a label by prefixing a name and a colon

IF Statements

- Simplest form:
IF <condition> THEN
 <statements(s)>
END IF;
- Add ELSE <statement(s)> if desired, as
IF . . . THEN . . . ELSE . . . END IF;
- Add additional cases by ELSEIF
<statements(s)>: IF ... THEN ... ELSEIF ...
THEN ... ELSEIF ... THEN ... ELSE ... END IF;

Example: IF

- Let's rate bars by how many customers they have, based on `Frequents(drinker,bar)`
 - <100 customers: 'unpopular'
 - 100-199 customers: 'average'
 - ≥ 200 customers: 'popular'
- Function `Rate(b)` rates bar b

Example: IF

```
CREATE FUNCTION Rate (IN b CHAR(20))  
  RETURNS CHAR(10) AS $$  
  DECLARE cust INTEGER;  
  BEGIN  
    cust := (SELECT COUNT(*) FROM Frequents  
             WHERE bar = b);  
    IF cust < 100 THEN RETURN 'unpopular';  
    ELSEIF cust < 200 THEN RETURN 'average';  
    ELSE RETURN 'popular';  
    END IF;  
  END;  
END;
```

Number of customers of bar b

IF statement

Loops

- Basic form:

```
<<<label>>>      LOOP  
    <statements>  
END LOOP;
```

- Exit from a loop by:

```
EXIT <label> WHEN <condition>
```

Example: Exiting a Loop

<<loop1>> LOOP

. . .

EXIT loop1 WHEN ...;

. . .

END LOOP;

← If this statement is executed and
the condition holds ...

← ... control winds up here

Other Loop Forms

- WHILE <condition> LOOP
 <statements>
END LOOP;
- Equivalent to the following LOOP:
LOOP
 EXIT WHEN NOT <condition>;
 <statements>
END LOOP;

Other Loop Forms

- FOR <name> IN <start> TO <end>
LOOP

 <statements>

END LOOP;

- Equivalent to the following block:

 <name> := <start>;

LOOP EXIT WHEN <name> > <end>;

 <statements>

 <name> := <name> + 1;

END LOOP;

Other Loop Forms

- FOR <name> IN REVERSE <start> TO <end> LOOP
 <statements>
END LOOP;
- Equivalent to the following block:
 <name> := <start>;
 LOOP EXIT WHEN <name> < <end>;
 <statements>
 <name> := <name> - 1;
END LOOP;

Other Loop Forms

- FOR <name> IN <start> TO <end>
BY <step> LOOP
 <statements>
END LOOP;
- Equivalent to the following block:
 <name> := <start>;
 LOOP EXIT WHEN <name> > <end>;
 <statements>
 <name> := <name> + <step>;
 END LOOP;

Queries

- General SELECT-FROM-WHERE queries are *not* permitted in PL/pgsql
- There are three ways to get the effect of a query:
 1. Queries producing one value can be the expression in an assignment
 2. Single-row SELECT ... INTO
 3. Cursors

Example: Assignment/Query

- Using local variable p and **Sells(bar, beer, price)**, we can get the price Cafe Chino charges for Odense Classic by:

```
p := (SELECT price FROM Sells  
      WHERE bar = 'C.Ch' AND  
            beer = 'Od.Cl.' );
```


SELECT ... INTO

- Another way to get the value of a query that returns one tuple is by placing **INTO** **<variable>** after the SELECT clause
- **Example:**

```
SELECT price INTO p FROM Sells
WHERE bar = 'C.Ch.' AND
      beer = 'Od.Cl.';
```

Cursors

- A *cursor* is essentially a tuple-variable that ranges over all tuples in the result of some query
- Declare a cursor *c* by:

```
DECLARE c CURSOR FOR <query>;
```

Opening and Closing Cursors

- To use cursor c , we must issue the command:

OPEN c ;

- The query of c is evaluated, and c is set to point to the first tuple of the result
- When finished with c , issue command:
CLOSE c ;

Fetching Tuples From a Cursor

- To get the next tuple from cursor c , issue command:

FETCH FROM c INTO x_1, x_2, \dots, x_n ;

- The x 's are a list of variables, one for each component of the tuples referred to by c
- c is moved automatically to the next tuple

Breaking Cursor Loops – (1)

- The usual way to use a cursor is to create a loop with a FETCH statement, and do something with each tuple fetched
- A tricky point is how we get out of the loop when the cursor has no more tuples to deliver

Breaking Cursor Loops – (2)

- Many operations return if a row has been found, changed, inserted, or deleted (SELECT INTO, UPDATE, INSERT, DELETE, FETCH)
- In plpgsql, we can get the value of the status in a variable called FOUND

Breaking Cursor Loops – (3)

- The structure of a cursor loop is thus:

```
<<cursorLoop>> LOOP
```

```
...
```

```
FETCH c INTO ... ;
```

```
IF NOT FOUND THEN EXIT cursorLoop;
```

```
END IF;
```

```
...
```

```
END LOOP;
```

Example: Cursor

- Let us write a procedure that examines `Sells(bar, beer, price)`, and raises by 10 the price of all beers at Cafe Chino that are under 30
- Yes, we could write this as a simple UPDATE, but the details are instructive anyway

Declarations

```
CREATE FUNCTION RaisePrices()
```

```
RETURNS VOID AS $$
```

```
DECLARE theBeer CHAR(20);  
        thePrice REAL;
```

Used to hold
beer-price pairs
when fetching
through cursor c

```
c CURSOR FOR
```

```
(SELECT beer, price FROM Sells  
WHERE bar = 'C.Ch.');
```

Returns Cafe Chino's
price list

Procedure Body

BEGIN

OPEN c;

<<menuLoop>> LOOP

FETCH c INTO theBeer, thePrice;

Check if the recent
FETCH failed to
get a tuple

EXIT menuLoop WHEN NOT FOUND;

IF thePrice < 30 THEN

UPDATE Sells SET price = thePrice + 10

WHERE bar = 'C.Ch.' AND beer = theBeer;

END IF;

END LOOP;

CLOSE c;

END;\$\$ LANGUAGE plpgsql;

If Cafe Chino charges less than
30 for the beer, raise its price at
at Cafe Chino by 10

Database-Connection Libraries

Host/SQL Interfaces Via Libraries

- The third approach to connecting databases to conventional languages is to use library calls
 1. C + CLI
 2. Java + JDBC
 3. Python + psycopg2

Three-Tier Architecture

- A common environment for using a database has three tiers of processors:
 1. *Web servers* – talk to the user.
 2. *Application servers* – execute the business logic
 3. *Database servers* – get what the app servers need from the database

Example: Amazon

- Database holds the information about products, customers, etc.
- Business logic includes things like “what do I do after someone clicks ‘checkout’ ?”
 - **Answer:** Show the “how will you pay for this?” screen

Environments, Connections, Queries

- The database is, in many DB-access languages, an *environment*
- Database servers maintain some number of *connections*, so app servers can ask queries or perform modifications
- The app server issues *statements*: queries and modifications, usually

JDBC

- *Java Database Connectivity* (JDBC) is a library for accessing a DBMS using Java as the host language
- >200 drivers available: PostgreSQL, MySQL, Oracle, ODBC, ...
- <http://jdbc.postgresql.org/>

Making a Connection

```
import java.sql.*;  
...  
Class.forName("org.postgresql.Driver");  
Connection myCon =  
    DriverManager.getConnection(...);  
...
```

The JDBC classes

Loaded by
forName

URL of the database
your name, and password
go here

The driver
for postgresql;
others exist

URL for PostgreSQL database

- `jdbc:postgresql://<host>[:<port>]/<database>?user=<user>&password=<password>`
- Alternatively use `getConnection` variant:
- `getConnection("jdbc:postgresql://<host>[:<port>]/<database>", <user>, <password>);`
- `DriverManager.getConnection("jdbc:postgresql://10.110.4.32:5434/postgres", "peter", "geheim");`

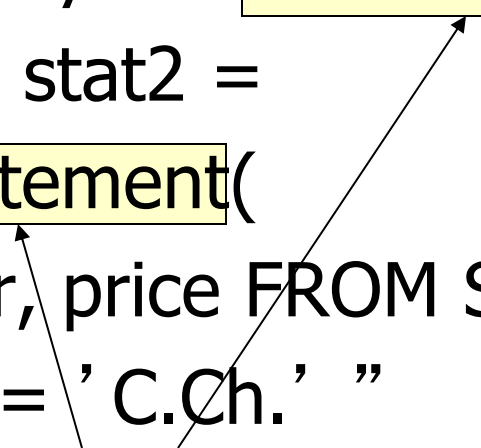
Statements

- JDBC provides two classes:
 1. *Statement* = an object that can accept a string that is a SQL statement and can execute such a string
 2. *PreparedStatement* = an object that has an associated SQL statement ready to execute

Creating Statements

- The Connection class has methods to create Statements and PreparedStatements

```
Statement stat1 = myCon.createStatement();  
PreparedStatement stat2 =  
    myCon.createStatement(  
        "SELECT beer, price FROM Sells " +  
        "WHERE bar = ' C.Ch.' "  
    );
```



`createStatement` with no argument returns a Statement; with one argument it returns a PreparedStatement

Executing SQL Statements

- JDBC distinguishes queries from modifications, which it calls “updates”
- Statement and PreparedStatement each have methods `executeQuery` and `executeUpdate`
 - For Statements: one argument – the query or modification to be executed
 - For PreparedStatements: no argument

Example: Update

- stat1 is a Statement
- We can use it to insert a tuple as:

```
stat1.executeUpdate(  
    "INSERT INTO Sells " +  
    "VALUES ('C.Ch.', 'Eventyr', 30)"  
);
```

Example: Query

- stat2 is a PreparedStatement holding the query "SELECT beer, price FROM Sells WHERE bar = ' C.Ch.' "
- **executeQuery** returns an object of class ResultSet – we'll examine it later
- The query:

```
ResultSet menu = stat2.executeQuery();
```

Accessing the ResultSet

- An object of type ResultSet is something like a cursor
- Method `next()` advances the “cursor” to the next tuple
 - The first time `next()` is applied, it gets the first tuple
 - If there are no more tuples, `next()` returns the value `false`

Accessing Components of Tuples

- When a ResultSet is referring to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet
- Method `getX(i)`, where *X* is some type, and *i* is the component number, returns the value of that component
 - The value must have type *X*

Example: Accessing Components

- Menu = ResultSet for query “SELECT beer, price FROM Sells WHERE bar = 'C.Ch.' ”
- Access beer and price from each tuple by:

```
while (menu.next()) {  
    theBeer = menu.getString(1);  
    thePrice = menu.getFloat(2);  
    /*something with theBeer and  
       thePrice*/  
}
```

Important Details

- Reusing a Statement object results in the ResultSet being closed
 - Always create new Statement objects using `createStatement()` or explicitly close ResultSets using the `close` method
- For transactions, for the Connection `con` use `con.setAutoCommit(false)` and explicitly `con.commit()` or `con.rollback()`
 - If `AutoCommit` is false and there is no commit: closing the connection = rollback

Python and Databases

- many different modules for accessing databases
- commercial: mxodbc, ...
- open source: pygresql, **psycopg2**, ...
- we use psycopg2
 - install using `pip install psycopg2`
 - import with `import psycopg2`

Connection String

- Database connection described by a connection string

- **Example:** `con_str = """`

`host=10.110.4.32`

`port=5434`

`dbname=postgres`

`user=peter`

`password=geheim`

`"""`

Making a Connection

- With the DB library imported and the connection string `con_str` available:

```
con = psycopg2.connect(con_str);
```

Function connect
in the DB API

Class is connection
because it is returned
by `psycopg2.connect(...)`

Cursors in Python

- Queries are executed for a cursor
- A cursor is obtained from connection
- Example:

```
cursor = con.cursor()
```

- Queries or modifications are executed using the `execute(...)` method
- Cursors can then be used in a `for`-loop

Example: Executing a Query

- Find all the bars that sell a beer given by the variable `beer`

```
beer = 'Od.Cl.'
```


```
cursor = con.cursor()
```

```
cursor.execute(
```

```
    "SELECT bar FROM Sells" +
```

```
    "WHERE beer = '%s' ;" % beer);
```

Remember this
variable is replaced
by the value of `beer`



Example: Tuple Cursors

```
bar = 'C.Ch.'  
cur = con.cursor()  
cur.execute("SELECT beer, price" +  
" FROM Sells" +  
" WHERE bar = " + bar + ";")  
for row in cur:  
    print row[0] + " for " + row[1]
```

Summary 12

More things you should know:

- Stored Procedures, PL/pgsql
- Declarations, Statements, Loops,
- Cursors, Tuple Variables
- Three-Tier Approach, JDBC, psycopg2